Towards Comparable Active Learning

April 25, 2023

1 Introduction

1.1 Contribution

2 Related Work

Version: Braindump

Many different algorithms have been proposed for active learning. In this work we focus on those approaches that have shown consistent results over the years as well as newer approaches that have demonstrated significant lifts in their initial experiments. AL algorithms can be categorized into two classes: Geometric approaches and uncertainty-based approaches. Geometric approaches include CoreSet [9] and TypiClust [3], which use clustering techniques to partition the data and then sample their unlabeled points based on the clusters. Uncertainty-based approaches include classic uncertainty sampling (based on Shannon-Entropy and the margin-score), BALD [5] and BADGE [1], which use metrics to measure the classifiers state.

Some previous work also aimed to provide a benchmark suite for active learning: The authors of [2] and [7] both focus on active learning in the image domain. While [2] discuss a new metric to measure AL performance, which they call "Label Efficiency" and provide experiments on many common configurations for data preparation, model training and other hyperparameters, [7] focuses on combined approaches of AL and semi-supervised learning to aid model training. The authors of [4] study models that are learned with AL techniques in the image and text domain. They test for several different properties of the models including robustness, response to compression techniques and final performance.

3 Overview

Table 1 shows a feature comparison between our proposed benchmark and several existing benchmarks in the literature, as well as methodological AL papers with an extensive experiments section.
TODO We include in this table any methodological paper that experiments on at least two domains.
TODO Define AL scenarios (really hard)

Paper	# Datasets	Domains	Scenarios	Oracle	RL Alg.
Beck et al. [2]	4	1	3	-	-
Hu et al. [4]	5	2	1	-	-
Li et al. [7]	5	1	1	-	-
Zhou et al. [11]	3	2	1	✓	-
Ours	5	2	2	✓	-

Table 1: Comparison of our benchmark with the existing literature

Manual NIPS Benchmark paper eval (out of 6):

- Has new method: 1

- Multiple domains: 2

- > 5 datasets: 3

- > 5 baselines: 4

3.1 Problem Description

Version: 1.0

Given two spaces $\mathcal{X} := \mathcal{R}^M$ and $\mathcal{Y} := \mathcal{R}^C$, a sample $\mathcal{D}_1, \ldots, \mathcal{D}_N \subseteq (\mathcal{X} \times \mathcal{Y})^*$ of sequences of pairs (x, y) from an unknown distribution p called datasets and a number $B \in \mathcal{N}$ with $B < |\mathcal{D}|$.

Given two functions $\ell : \mathcal{Y} \times \mathcal{Y} \to \mathcal{R}$ called loss, and $A : (\mathcal{X} \times \mathcal{Y})^* \times \mathcal{X}^* \to \mathcal{Y}^{\mathcal{X}}$ called learning algorithm, find a function

$$a : (\mathcal{X} \times \mathcal{Y})^* \times \mathcal{X}^* \to \{0, 1\}^* \qquad \text{\scriptsize (equivariant in the second argument)}$$

called acquisition function, s.t. the expected loss of a model learned on all predictors plus B acquired targets is minimal:

$$\mathbb{E}_{\mathcal{D}\text{train}, \mathcal{D}\text{test} \sim p} \, \text{avg}_{(x,y) \in \mathcal{D}\text{test}} \, \ell(y, \hat{y}(x))$$
$$\text{with } \hat{y} := A((\mathcal{D}\text{train}_{n_1}, \ldots, \mathcal{D}\text{train}_{n_B}), \mathcal{D}\text{train}|_{\mathcal{X}})$$
$$n_b := \text{index}(a((\mathcal{D}\text{train}_{n_1}, \ldots, \mathcal{D}\text{train}_{n_{b-1}}), \mathcal{D}\text{train}|_{\mathcal{X}})), \quad b \in 1{:}B$$

(*We would need to switch from lowest expected loss to highest AUC*)

Since combinatorial problem of finding the optimal subset $\mathcal{D}_{\text{train}}$ is computational not feasible, we apply two relaxations to the problem. Firstly, we allow sequential construction of $\mathcal{D}_{\text{train}}$ and secondly, we limit the choice of data points in every iteration to a random sample of size τ.

To construct the active learning setting, we suppress the labels $\mathcal{Y}$ of $\mathcal{D}_{\text{train}}$ to form the unlabeled pool $\mathcal{U} := \mathcal{X}^*$ and form and initial labeled pool by uniformly sampling k number of instances per class from $\mathcal{U}$ and recovering their label $\mathcal{L} := (\mathcal{X}, \mathcal{Y})^{k*C}$.

Following [11], the quality of an active learning algorithm is evaluated by an "anytime" protocol that incorporates classification performance at every iteration, not just the final performance after the budget is exhausted. We employ the normalized area under the accuracy curve (AUC):

$$\text{auc}(\mathcal{U}, \mathcal{L}, \hat{y}, B) := \frac{1}{B} \sum_{i=1}^{B} \text{Acc}(y_{test}, \hat{y}_i(x_{test})) \tag{1}$$

Where $\hat{y}_i$ is the retrained classification model after the i-*th* instance was selected.

Framing AL as RL

We define the active learning process as an adapted reinforcement learning loop $(S, A, \tau, \Omega, \omega)$ where an environment iteratively will expose a state $s \in S$ to an agent Ω, which will choose actions $a \in A$. For each iteration i the environment samples a subset of size τ of unlabeled instances $u^{(i)} \underset{\tau}{\sim} \mathcal{U}^{(i)}$, constructs the state $s^{(i)} := \omega(u^{(i)})$ and presents it to the agent to select an action $a^{(i)} := \Omega(s^{(i)})$. The action $a^{(i)}$ is the index of the selected instance in $u^{(i)}$ out of all possible indices $A := [1 \ldots \tau]$. This process is repeated B times $i := [1 \ldots B]$.

Algorithm 1 Active Learning OLD

Require: $\mathcal{U}$ ▷ Unlabeled Pool
Require: τ ▷ Unlabeled Sample Size
Require: Ω ▷ AL Agent
Require: ω ▷ Environment State function
1: $\mathcal{L}^{(1)} \leftarrow \text{seed}(\mathcal{U})$ ▷ Create the initial labeled set
2: $\mathcal{U}^{(1)} \leftarrow \mathcal{U}$
3: **for** $i := 1 \ldots B$ **do**
4:    $\text{acc}^{(i)} \leftarrow \text{Retrain}(\mathcal{L}^{(i)})$ ▷ Retrain($\mathcal{L}^{(i)}$) is shorthand for Retrain($\mathcal{L}^{(i)}, \mathcal{L}^{\text{test}}, \hat{y}_\theta, e^{\max}$)
5:    $u^{(i)} \underset{\tau}{\sim} \mathcal{U}^{(i)}$
6:    $s^{(i)} \leftarrow \omega(u^{(i)})$
7:    $a^{(i)} \leftarrow \Omega(s^{(i)})$ ▷ $a^{(i)}$ is an index inside of $u^{(i)}$
8:    $y^{(i)} \leftarrow \text{label}(u_a^{(i)})$ ▷ $u_a^{(i)}$ is shorthand for $u_{a^{(i)}}^{(i)}$
9:    $\mathcal{L}^{(i+1)} \leftarrow \mathcal{L}^{(i)} \cup \{(u_a^{(i)}, y^{(i)})\}$
10:    $\mathcal{U}^{(i+1)} \leftarrow \mathcal{U}^{(i)} \setminus \{u_a^{(i)}\}$
11: **end for**
12: **return** $\frac{1}{B} \sum_{i=1}^{B} \text{acc}^{(i)}$

Algorithm 2 Retrain

Require: $\mathcal{L}$ ▷ Labeled Pool
Require: $\mathcal{L}^{\text{test}}$ ▷ Labeled Test Data
Require: $\hat{y}_\theta$ ▷ Classification Model
Require: $e^{\max}$ ▷ Maximum Epochs
1: $\text{loss}^* \leftarrow \infty$
2: **for** $i := 1 \ldots e^{\max}$ **do**
3:    $\theta_{i+1} \leftarrow \theta_i - \eta \nabla_\theta \ell(\mathcal{L}, \hat{y}_\theta)$
4:    $\text{loss}_i \leftarrow \ell(\mathcal{L}^{\text{test}}, \hat{y}_\theta)$
5:    **if** $\text{loss}_i < \text{loss}^*$ **then**
6:      $\text{loss}^* \leftarrow \text{loss}_i$
7:    **else**
8:      Break
9:    **end if**
10: **end for**
11: **return** $\text{Acc}(\mathcal{L}^{\text{test}}, \hat{y}_\theta)$

Algorithm 3 Oracle OLD

Require: $\mathcal{U}$ ▷ Unlabeled Pool
Require: τ ▷ Unlabeled Sample Size
Require: Ω ▷ AL Agent
Require: ω ▷ Environment State function
1: $\mathcal{L}^{(1)} \leftarrow \text{seed}(\mathcal{U})$ ▷ Create the initial labeled set
2: $\mathcal{U}^{(1)} \leftarrow \mathcal{U}$
3: **for** $i := 1 \ldots B$ **do**
4:    $\text{acc}^{(i)} \leftarrow \text{Retrain}(\mathcal{L}^{(i)})$ ▷ Retrain($\mathcal{L}^{(i)}$) is shorthand for Retrain($\mathcal{L}^{(i)}, \mathcal{L}^{\text{test}}, \hat{y}_\theta, e^{\max}$)
5:    $u^{(i)} \underset{\tau}{\sim} \mathcal{U}^{(i)}$
6:    $r* \leftarrow -\infty$
7:    $j* \leftarrow -1$
8:    **for** $j := 1 \ldots \tau$ **do**   ▷ Testing every unlabeled point
9:      $y^{(j)} \leftarrow \text{label}(u_j^{(i)})$
10:      $\mathcal{L}^{(j)} \leftarrow \mathcal{L}^{(i)} \cup \{(u_j^{(i)}, y^{(j)})\}$
11:      $\text{acc}^{(j)} \leftarrow \text{Retrain}(\mathcal{L}^{(j)})$
12:      $r^{(j)} \leftarrow \text{acc}^{(j)} - \text{acc}^{(i)}$
13:      **if** $r^{(j)} > r^*$ **then**   ▷ Select point with largest increase in performance
14:        $r* \leftarrow r^{(j)}$
15:        $j* \leftarrow j$
16:      **end if**
17:    **end for**
18:    $y^{(i)} \leftarrow \text{label}(u_{j*}^{(i)})$
19:    $\mathcal{L}^{(i+1)} \leftarrow \mathcal{L}^{(i)} \cup \{(u_{j*}^{(i)}, y^{(i)})\}$
20:    $\mathcal{U}^{(i+1)} \leftarrow \mathcal{U}^{(i)} \setminus \{u_{j*}^{(i)}\}$
21: **end for**
22: **return** $\frac{1}{B} \sum_{i=1}^{B} \text{acc}^{(i)}$

3.2 Datasets

In this work we use 4 different datasets from two domains. For vector data, we use DNA and Splice (Source: [8]) and for image data we use FashionMNIST [10] and Cifar10 [6].

We would like to point out that these datasets can be considered "toy-datasets" and therefore not relevant for practical purposes. This might be true if we aimed to develop novel classification models on these dataset, however we are solely focused on comparing different AL algorithms in this paper. Our core assumption is that a well-performing algorithm in our benchmark will transfer well into more practical use-cases.

Adapting the experimental setting from [3] we offer all our datasets in the raw setting as well as pre-encoded by a fixed embedding model that was trained in an unsupervised manner. The pre-encoded datasets offer multiple advantages in that they require smaller classification models (i.e. small MLPs) and in general much less sampled datapoints to archive close to upper bound performance (measured with the full dataset). They unify experimental setups for different datasets and therefore make them much more comparable.

4 Methodology

4.1 Reproducibility

A big focus in this work is to provide an experimental setup that is fully reproducible independent from the dataset, classification model or AL algorithm used. In our opinion, an evaluation on one dataset with a given seed should always be done on the same sequence of batches $x_0, \ldots, x_B$. Even though different AL algorithms will pick different samples from these batches, making them unavailable for sampling in later batches, the theoretical decision tree for every possible choice in every iteration i should stay the same. This behavior is not possible with the default choice of setting a global seed at the start of the experiment, since a single additional random draw from the random number generator completely changes the decision tree for the batches. This additional random number might be drawn during the initialization of the classification model, or the AL algorithm, or even during every AL iteration if Ω is stochastic. This problem also applies to the initialization of the classification models θ, the initialization and querying of $Omega$ and the drawn mini batches during the training of θ. The desired control over all these processes can be archived by assigning a separate seeded random number generator to all these processes. To the best of our knowledge, we are the first work that discusses this issue and proposes a solution for it. We hypothesize that the insufficient setup with global seeds contributes to the on-going problem of inconsistent results of AL algorithms in different papers.

4.2 Comparibility

We split the evaluation model that generates accuracy scores and the model that generates features for the AL agent.
This allows for dropout / ensemble approaches to invest more compute to generate features
The normal evaluation model is available as a fallback of course
This way, the comparison is always fair

4.3 Evaluation

We compare different AL algorithms based on their archived AUC score (Eq. 1). This score does incorporate performance in early stages (low budget) as well as capabilities to push the classifier in later stages (high

Dataset	Architecture	Optimizer	LR	Weight Decay	Dropout
Splice	[24, 12]	NAdam	0.0012	5.9e-5	0
DNA	[24, 12]	NAdam	0.0398	3.6e-5	0
USPS	[24, 12]	Adam	0.0081	1.5e-6	0
FashionMnist	linear	-	-	-	-
Cifar10	linear	-	-	-	-

Table 2: Classifier architectures and optimized hyperparameters per dataset. Numbers in brackets signify a MLP with corresponding hidden layers.

budget). A good AL algorithm should be able to perform well in both scenarios.
TODO

5 Implementation Details

5.1 State Space

Version: 0.1 Remove analogies to RL
Since every AL agent needs a different state space our environment exposes a full state to the agent, so that the agent has full control of what information will be used.

The state can include the following information:

- The entire labeled dataset $\mathcal{L}^{(i)}$

- The entire unlabeled dataset $\mathcal{U}^{(i)}$

- A histogram of labeled points per class (count)

- The available budget B

- Number of added datapoints $|\mathcal{L}| - |\mathcal{S}|$

- The initial validation accuracy $\text{acc}^{(0)}$ and current validation accuracy $\text{acc}^{(i)}$

- The current classification model including all model weights $\theta^{(i)}$

- The current optimizer including it's full state

- The current sample of unlabeled points $u^{(i)}$

5.2 Choosing the Classifier

Traditionally, the classifier is chosen per dataset so that it is capable of solving the dataset close to the SOTA performance reported in the literature. Similar to our hypothesis in section 4.3 we hypothesize that AL algorithms will perform similarly on small classifiers and more complex ones, so that the overall ranking of algorithms stays the same. TODO: Transform this into full hypothesis incl. formatting?
On the basis of this hypothesis we opt to use smaller classifiers that still solve the dataset to a reasonable degree. Smaller classifiers also require fewer labeled datapoints to acquire performance close to the upper bound. This behavior can be observed in Figure TODO

For every dataset the chosen architecture's hyperparameters are optimized by to archive maximum upper bound performance. For an overview of architectures and hyperparameters please refer to Table 2.

	Splice	DNA	USPS
Oracle	0.835 +- 0.01	0.879 +- 0.01	0.870 +- 0.01
SAL_23_01_16	0.808 +- 0.01	0.857 +- 0.02	0.865 +- 0.01
Coreset_Greedy	0.808 +- 0.01	0.853 +- 0.02	0.849 +- 0.01
MarginScore	0.806 +- 0.01	0.855 +- 0.02	0.864 +- 0.01
ShannonEntropy	0.808 +- 0.01	0.855 +- 0.02	0.864 +- 0.01
RandomAgent	0.797 +- 0.01	0.823 +- 0.02	0.831 +- 0.01

	Cifar10	FashionMnist
Oracle	0.718 +- 0.01	0.721 +- 0.01
SAL_23_01_16	0.660 +- 0.01	0.667 +- 0.01
Coreset_Greedy	0.668 +- 0.01	0.674 +- 0.01
MarginScore	0.674 +- 0.01	0.677 +- 0.01
ShannonEntropy	0.671 +- 0.01	0.680 +- 0.01
RandomAgent	0.632 +- 0.02	0.633 +- 0.01

5.3   Training the Classifier

Version: 0.1 Full rework needed

Generally, the classification model can be trained in two ways. Either you reset the parameters after each AL iteration and train the classifier from scratch with the updated labeled set $\mathcal{L}^{(i)}$, or you retain the previous state and only fine-tune the classifier on $\mathcal{L}^{(i)}$ for a reduced number of epochs. In this work we use the fine-tuning method for raw datasets to save computation, while we use the from-scratch training for embedded dataset, since they have very small classifiers and this method generally produces better results. Our fine-tuning scheme always trains for at least one epoch and employs an aggressive early stopping after that. The early stopping has patience 0, so it will stop as soon as the validation loss does no longer decrease.

6 Experiments

6.1 Splice Embedded

Figure 1: Results for all algorithms on the pre-encoded Splice dataset

Figure 2: Critical Difference Diagram for Splice where every restart of the algorithm is one sample for the Wilcoxon-Holm method

Figure 3: Critical Difference Diagram for Splice, DNA, USPS and Cifar10

7 Ablation Studies

- Setting τ to $|\mathcal{U}|$
- Reduction of the test set for speed

References

[1] Jordan T Ash, Chicheng Zhang, Akshay Krishnamurthy, John Langford, and Alekh Agarwal. Deep batch active learning by diverse, uncertain gradient lower bounds. In *International Conference on Learning Representations*, 2020.

[2] Nathan Beck, Durga Sivasubramanian, Apurva Dani, Ganesh Ramakrishnan, and Rishabh Iyer. Effective evaluation of deep active learning on image classification tasks. *arXiv preprint arXiv:2106.15324*, 2021.

[3] Guy Hacohen, Avihu Dekel, and Daphna Weinshall. Active learning on a budget: Opposite strategies suit high and low budgets. *arXiv preprint arXiv:2202.02794*, 2022.

[4] Qiang Hu, Yuejun Guo, Maxime Cordy, Xiaofei Xie, Wei Ma, Mike Papadakis, and Yves Le Traon. Towards exploring the limitations of active learning: An empirical study. In *2021 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 917–929. IEEE, 2021.

[5] Andreas Kirsch, Joost Van Amersfoort, and Yarin Gal. Batchbald: Efficient and diverse batch acquisition for deep bayesian active learning. *Advances in neural information processing systems*, 32, 2019.

[6] Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. 2009.

[7] Yu Li, Muxi Chen, Yannan Liu, Daojing He, and Qiang Xu. An empirical study on the efficacy of deep active learning for image classification. *arXiv preprint arXiv:2212.03088*, 2022.

[8] Information Engineering Graduate Institute of Taiwan University. Libsvmtools.

[9] Ozan Sener and Silvio Savarese. Active learning for convolutional neural networks: A core-set approach. *arXiv preprint arXiv:1708.00489*, 2017.

[10] Han Xiao, Kashif Rasul, and Roland Vollgraf. Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*, 2017.

[11] Yilun Zhou, Adithya Renduchintala, Xian Li, Sida Wang, Yashar Mehdad, and Asish Ghoshal. Towards understanding the behaviors of optimal deep active learning algorithms. In *International Conference on Artificial Intelligence and Statistics*, pages 1486–1494. PMLR, 2021.

A Comparison of different sample sizes τ

Splice

random auc 0.815
entropy tau=N auc 0.82
entropy tau=200 auc 0.818
entropy tau=20 auc 0.829
upper bound
98% Percentile