

Figure 1: AUC vs. # of added datapoints. The graph shows the performance of various active learning methods as more data points are added. The Oracle method (red) consistently achieves the highest AUC, while the Random method (grey) serves as a baseline. The other methods (margin, Entropy, coreset, TypiClust, badge, BALD) show varying degrees of improvement over the Random baseline, with TypiClust and badge performing relatively well among the non-Oracle methods.

Method	AUC
Oracle	0.501
Random	0.45
margin	0.45
Entropy	0.449
coreset	0.453
TypiClust	0.437
badge	0.454
BALD	0.435

